

Partições



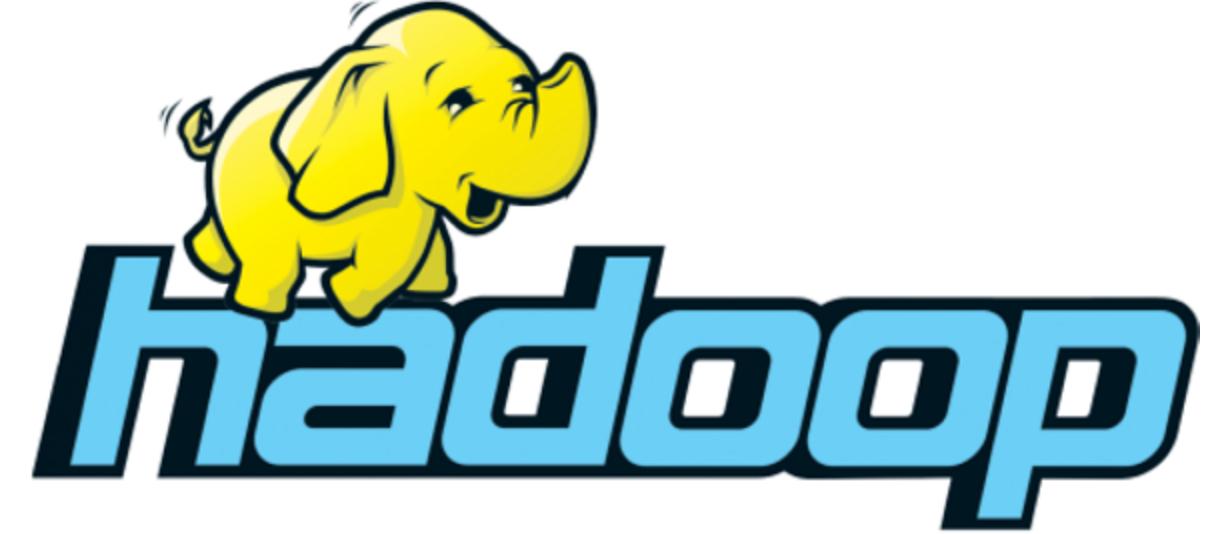
Bucketing



Compressão



Windowing Functions



Partições



Bucketing

Otimização de Performance



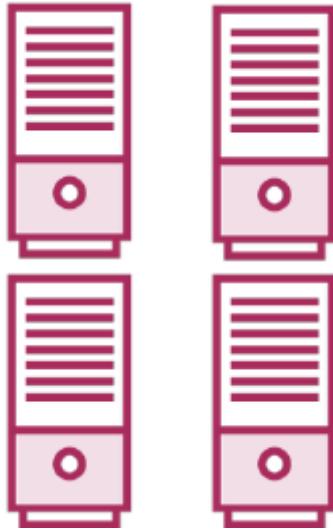
Compressão



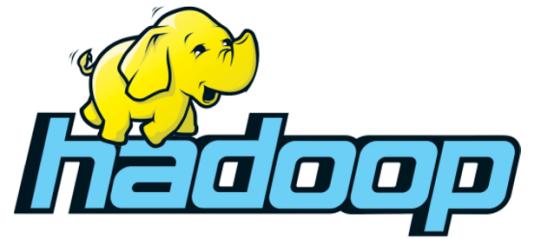
Windowing Functions

Boas práticas

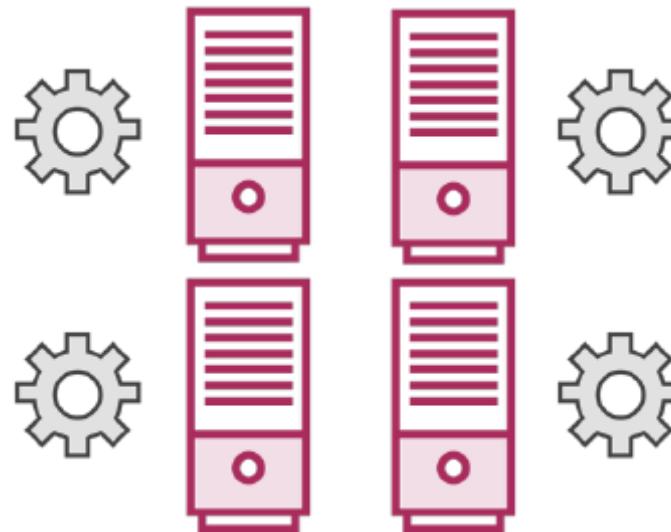
DataWarehouse



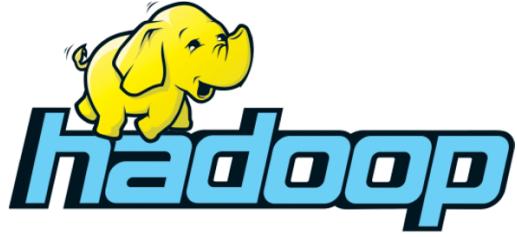
Dados são armazenados em múltiplas máquinas do cluster.



Framework Distribuído



Fundação Apache – Open Source



HDFS

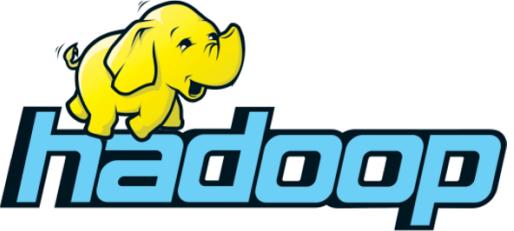
MapReduce

YARN

**Sistema
de
arquivos**

**Paradigma
de
computação
paralela**

Gerenciador



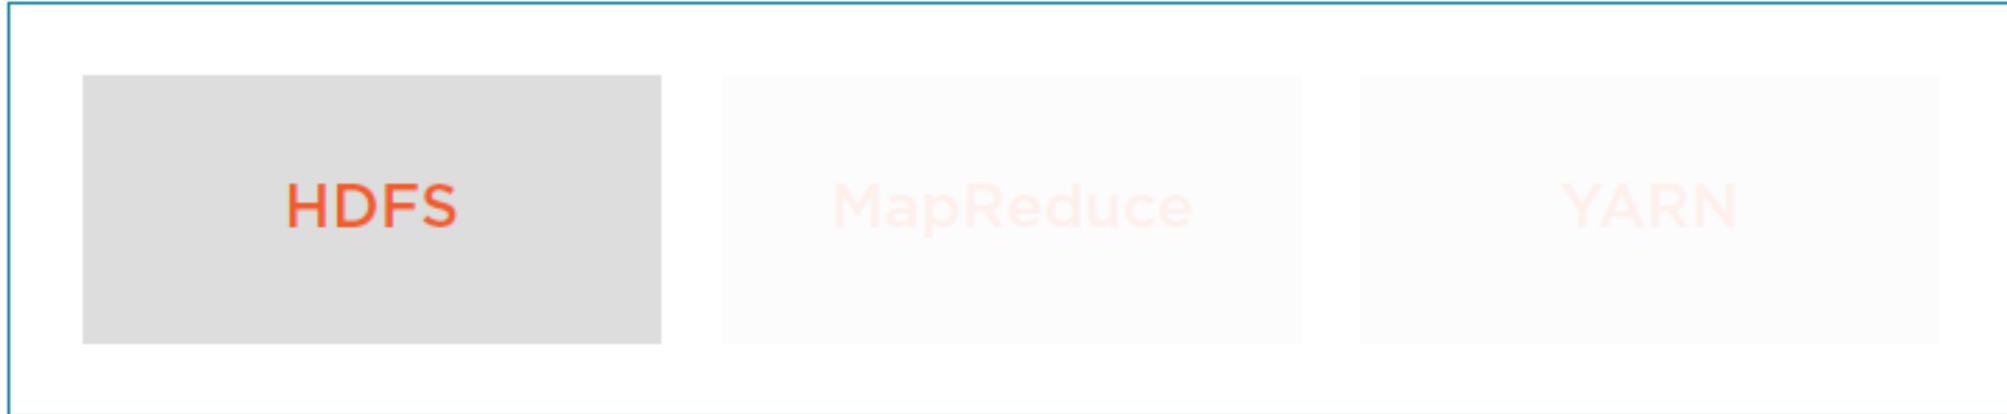
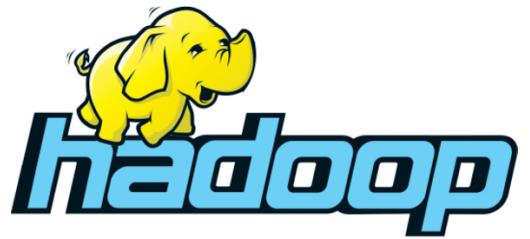
HIVE

HDFS

MapReduce

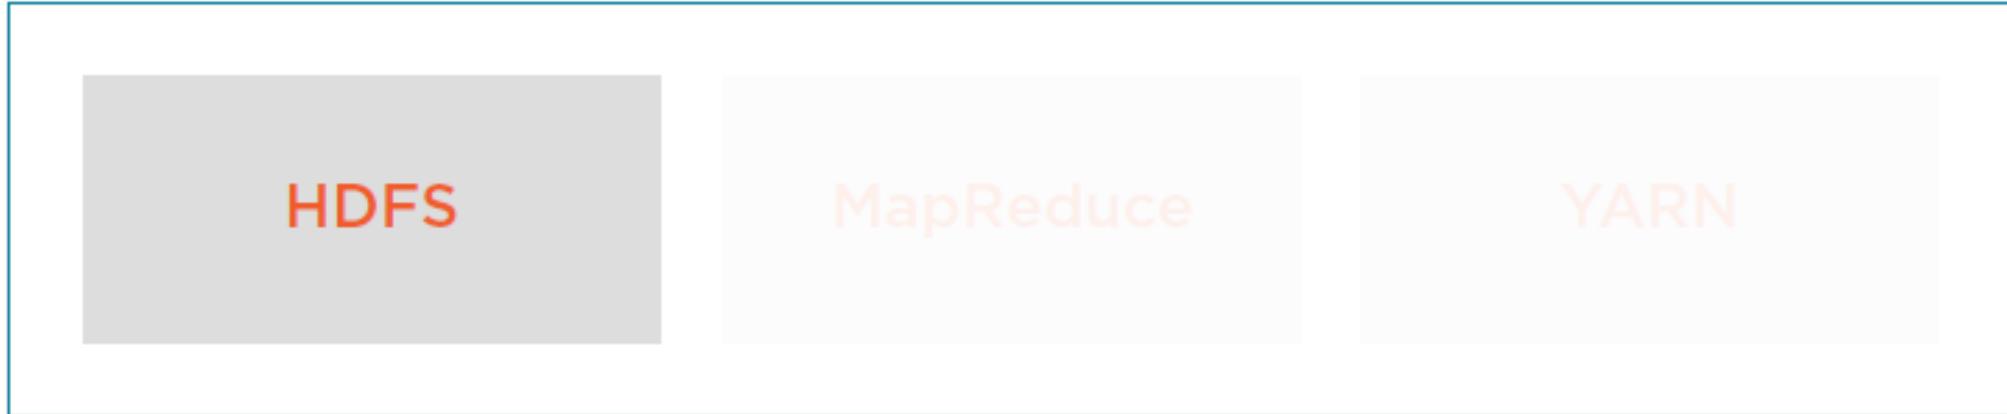
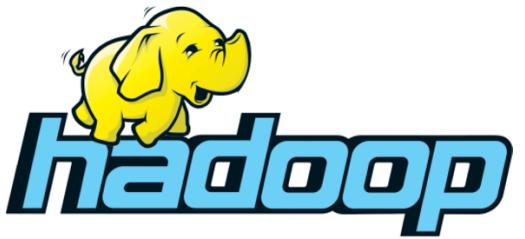
YARN

HIVE Query Language



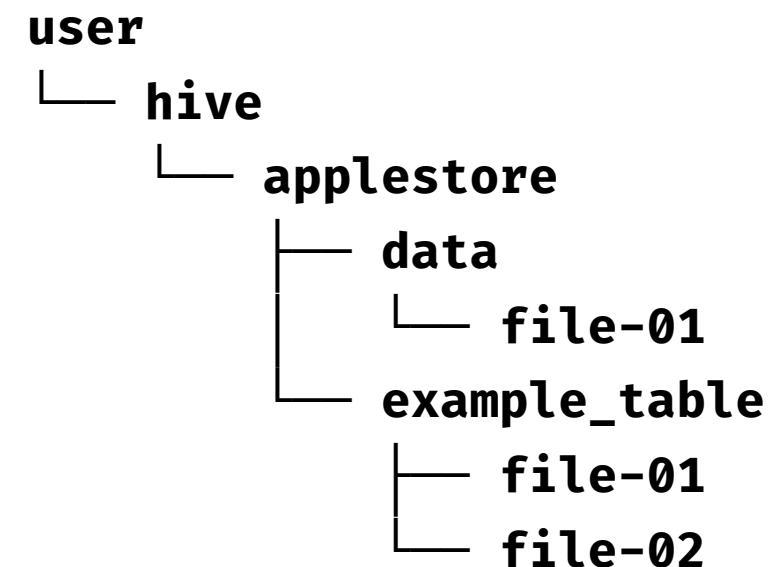
Namenodes
Datanodes
Replicação
Criação

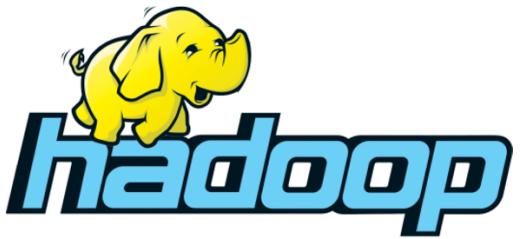
```
user
└── hive
    └── applestore
        ├── data
        │   └── file-01
        └── example_table
            ├── file-01
            └── file-02
```



Access via:

/user/hive/applestore/data/file-01

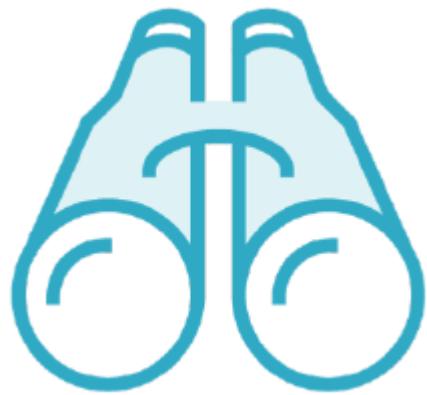




user
└── **hive**
 └── **applestore**
 ├── **data**
 │ └── **file-01**
 └── **example_table**
 ├── **file-01**
 └── **file-02**



Partições



São usadas para acelerar cláusula
WHERE

Partições são **PASTAS** no HDFS



Partições

pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transac	cat_estado
1	SANTADER	K	1000	599	SP
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP
5	RAPPI	2	500000	199	RR
6	NUBANK	K	25000	88	ES
7	GOOGLE	D	400000	45	ES
8	FACEBOOK	D	500000	25	RR

SELECT * FROM table WHERE CAT_ESTADO=“SP”



Partições



SP



RJ

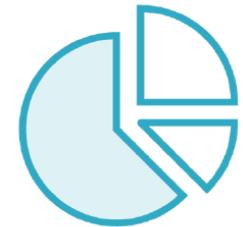
Red vertical line separating the first two columns from the second two columns.

PR

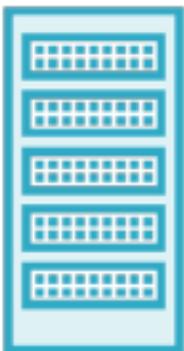
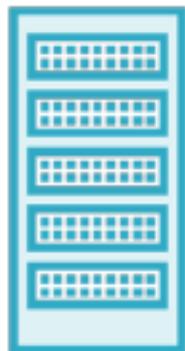
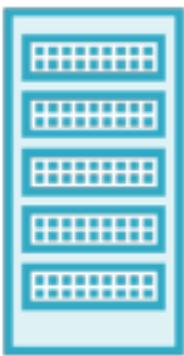
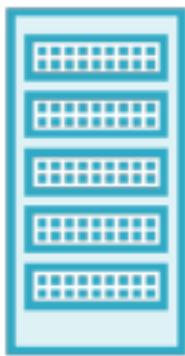


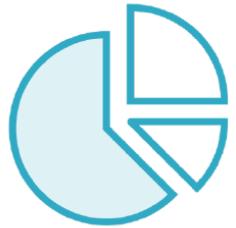
ES





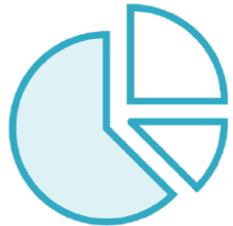
Partições





Partições

/user/hive/applestore/data/



Partições

/user/hive/applestore/data/

/cat_estado=SP



/cat_estado=RJ



/cat_estado=PR



/cat_estado=ES





Partições

/user/hive/applestore/data/

/cat_estado=SP

/file-01



Partições

/user/hive/applestore/data/cat_estado=SP/file-01

pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transacao	cat_estado
1	SANTADER	K	1000	599	SP
4	IFOOD	D	100000	99	SP

cat_estado=SP

pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transacao	cat_estado
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ

cat_estado=RJ

pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transacao	cat_estado
5	RAPPI	2	500000	199	PR
8	FACEBOOK	D	500000	25	PR

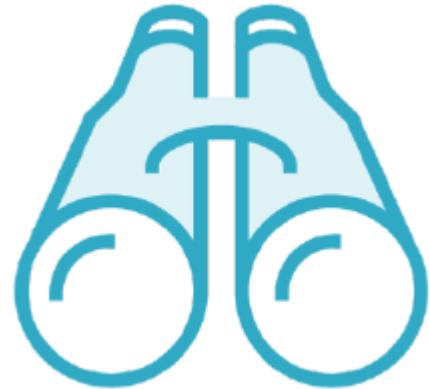
cat_estado=PR

pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transacao	cat_estado
6	NUBANK	K	25000	88	ES
7	GOOGLE	D	400000	45	ES

cat_estado=ES



Partições



- Devem ser aplicadas à coluna que você queira **filtrar** mais frequentemente
- É possível fazer mais de uma partição
- Balanceamento de carga
- Milhares de partições podem trazer grande overhead



Compressão



Compressão

Row Storage

1	SANTADER	K	1000	599	SP
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP

Column Storage

pk_cd_cnpj	cat_segmento	vlr_transacao	cat_estado
1	K	1000	SP
2	K	25000	RJ
3	2	400000	RJ
4	D	100000	SP
5	2	500000	RR
6	K	25000	ES
7	D	400000	ES
8	D	500000	RR



Compressão

Row Storage

1	SANTADER	K	1000	599	SP
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP

Column Storage

pk_cd_cnpj	cat_segmento	vlr_transacao	cat_estado
1	K	1000	SP
2	K	25000	RJ
3	2	400000	RJ
4	D	100000	SP
5	2	500000	PR
6	K	25000	ES
7	D	400000	ES
8	D	500000	PR

Mais fácil
de
comprimir!

vlr_transacao
0
...
0
1000
NaN
...
NaN
NaN

→ 500x0-1x1000-400xNaN



Compressão

Row Storage

1	SANTADER	K	1000	599	SP
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP

Column Storage

pk_cd_cnpj	cat_segmento	vlr_transacao	cat_estado
1	K	1000	SP
2	K	25000	RJ
3	2	400000	RJ
4	D	100000	SP
5	2	500000	RR
6	K	25000	BS
7	D	400000	BS
8	D	500000	RR

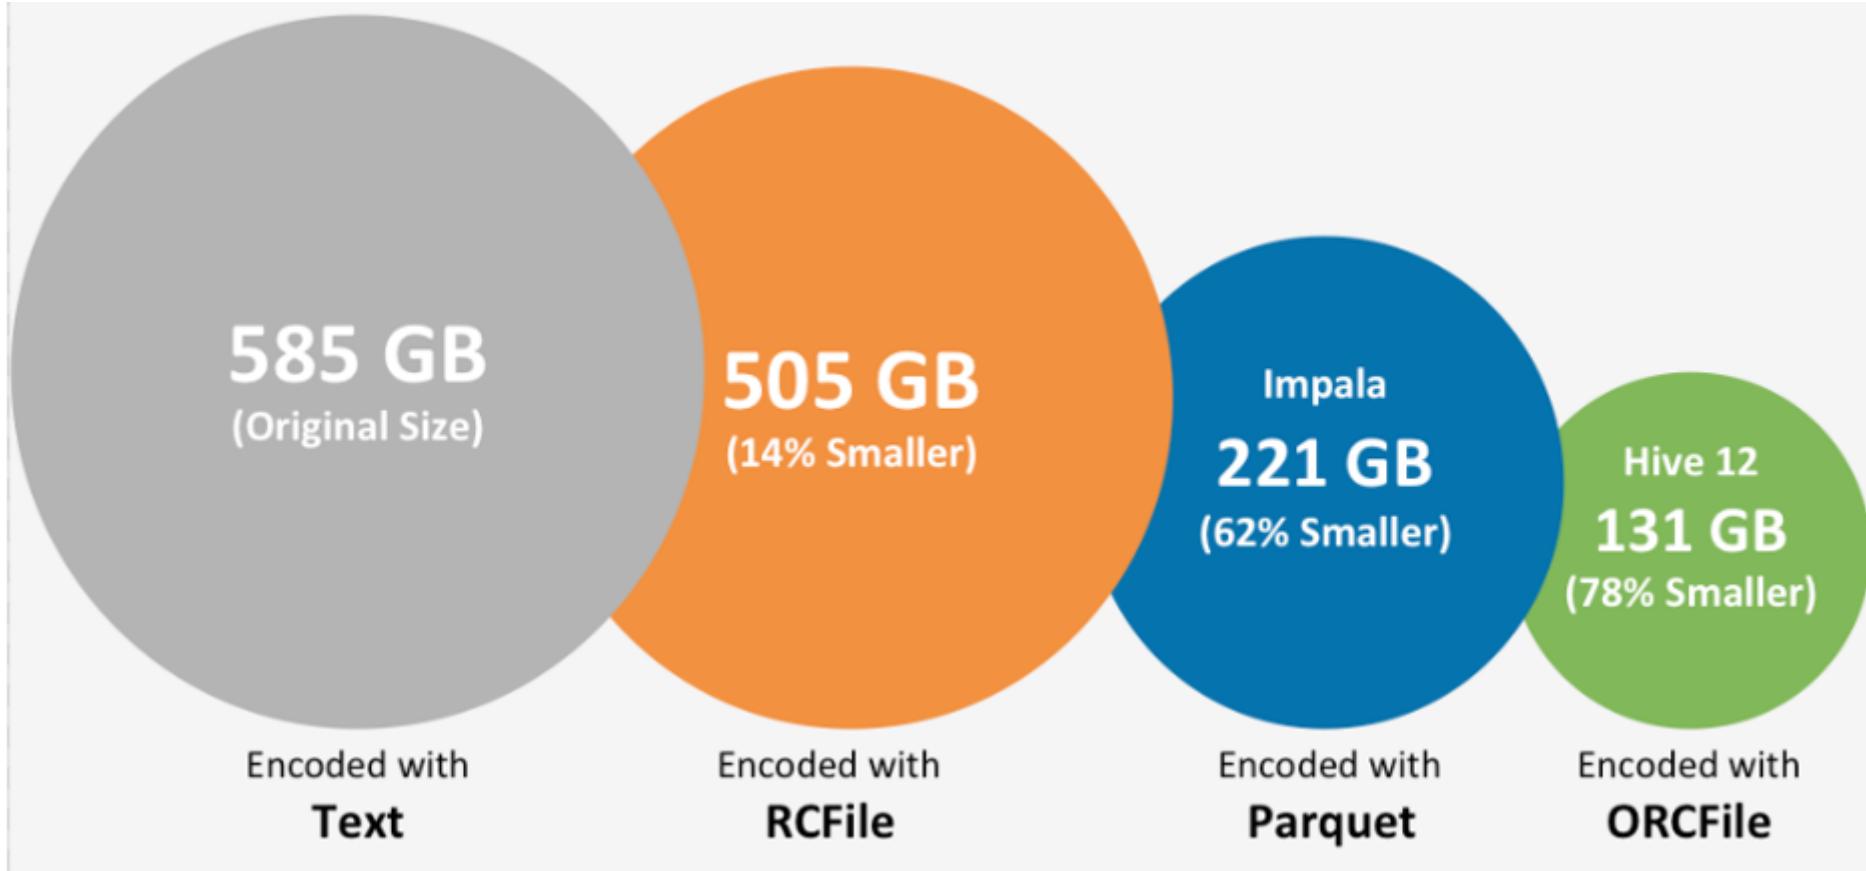
Mais fácil
de
comprimir!

fl_target
0
1
1
0
0
0
1
0



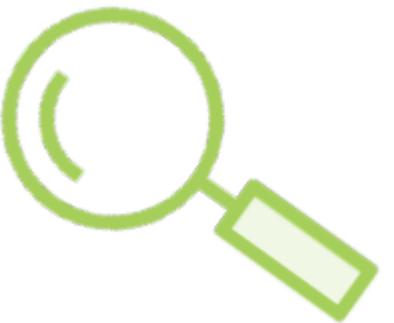
Compressão

Comparação





Bucketing





Bucketing



São usadas para acelerar cláusula
JOINS

Partições são **ARQUIVOS** no HDFS

Bucketing é o equivalente lógico de
Hash



Bucketing





Bucketing





Bucketing



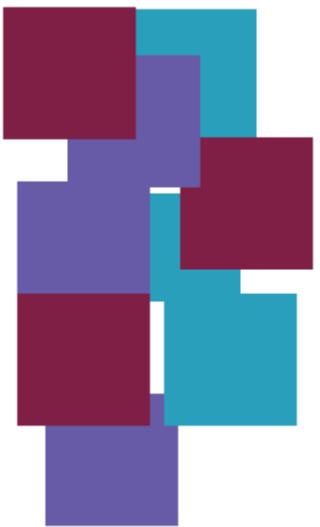


Bucketing





Bucketing

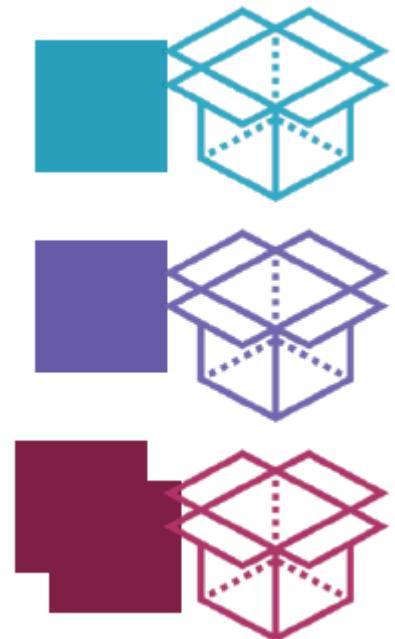




Bucketing



—





Bucketing

pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transac	cat_estado
1	SANTADER	K	1000	599	SP
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP
5	RAPPI	2	500000	199	RR
6	NUBANK	K	25000	88	ES
7	GOOGLE	D	400000	45	ES
8	FACEBOOK	D	500000	25	RR

3 Buckets – pk_cd_cnpj

Para inteiros, a função de hash é simples (módulo %)



Bucketing

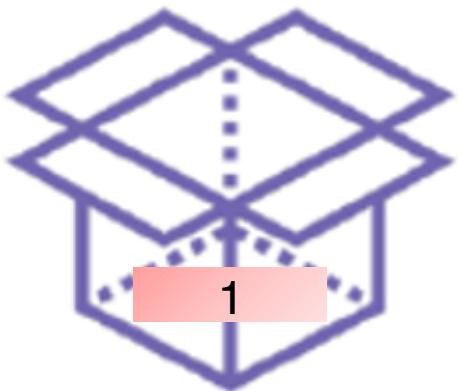
pk	cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transac	cat_estado
1		SANTADER	K	1000	599	SP
2		UBER	K	25000	58	RJ
3		ITAU	2	400000	66	RJ
4		IIFOOD	D	100000	99	SP
5		RAPPI	2	500000	199	RR
6		NUBANK	K	25000	88	RS
7		GOOGLE	D	400000	45	RS
8		FACEBOOK	D	500000	25	RR

$$1 \% 3 = 1$$

0



1



2





Bucketing

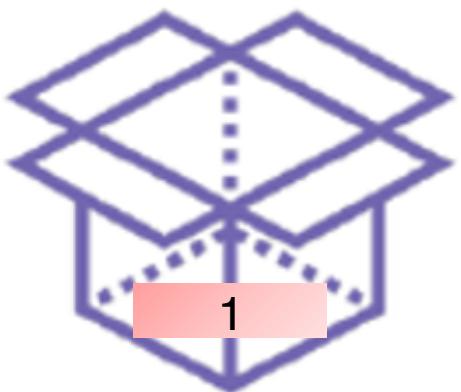
pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transac	cat_estado
1	SANTADER	K	1000	599	SP
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP
5	RAPPI	2	500000	199	RR
6	NUBANK	K	25000	88	ES
7	GOOGLE	D	400000	45	ES
8	FACEBOOK	D	500000	25	RR

2 % 3 = 2

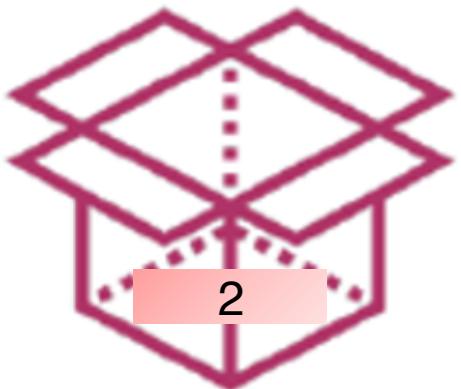
0



1



2



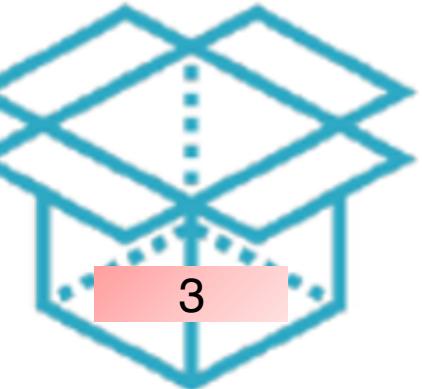


Bucketing

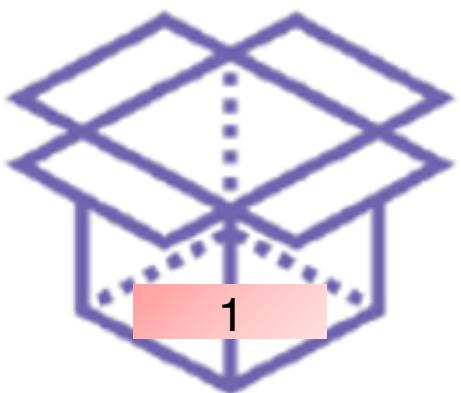
pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transac	cat_estado
1	SANTADER	K	1000	599	SP
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP
5	RAPPI	2	500000	199	RR
6	NUBANK	K	25000	88	ES
7	GOOGLE	D	400000	45	ES
8	FACEBOOK	D	500000	25	RR

3 % 3 = 0

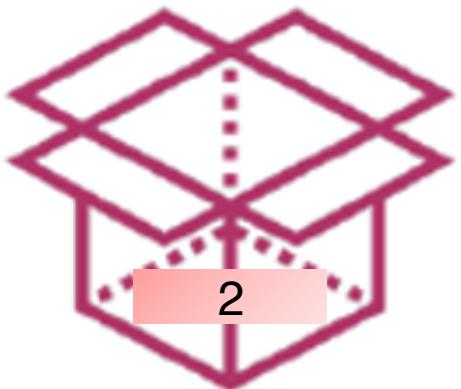
0



1



2





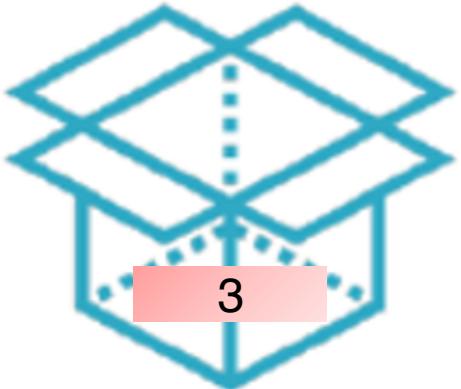
Bucketing

pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transac	cat_estado
1	SANTADER	K	1000	599	SP
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP
5	RAPPI	2	500000	199	RR
6	NUBANK	K	25000	88	ES
7	GOOGLE	D	400000	45	ES
8	FACEBOOK	D	500000	25	RR

4 % 3 =

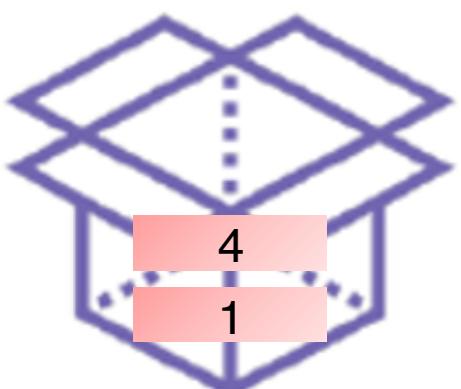
1

0

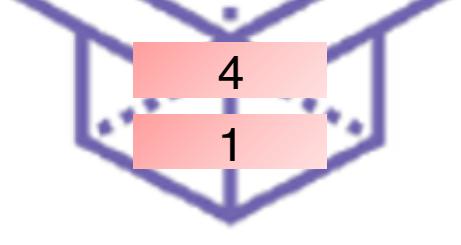


3

1

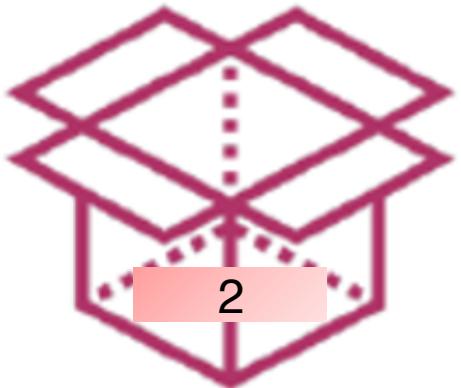


4



1

2



2

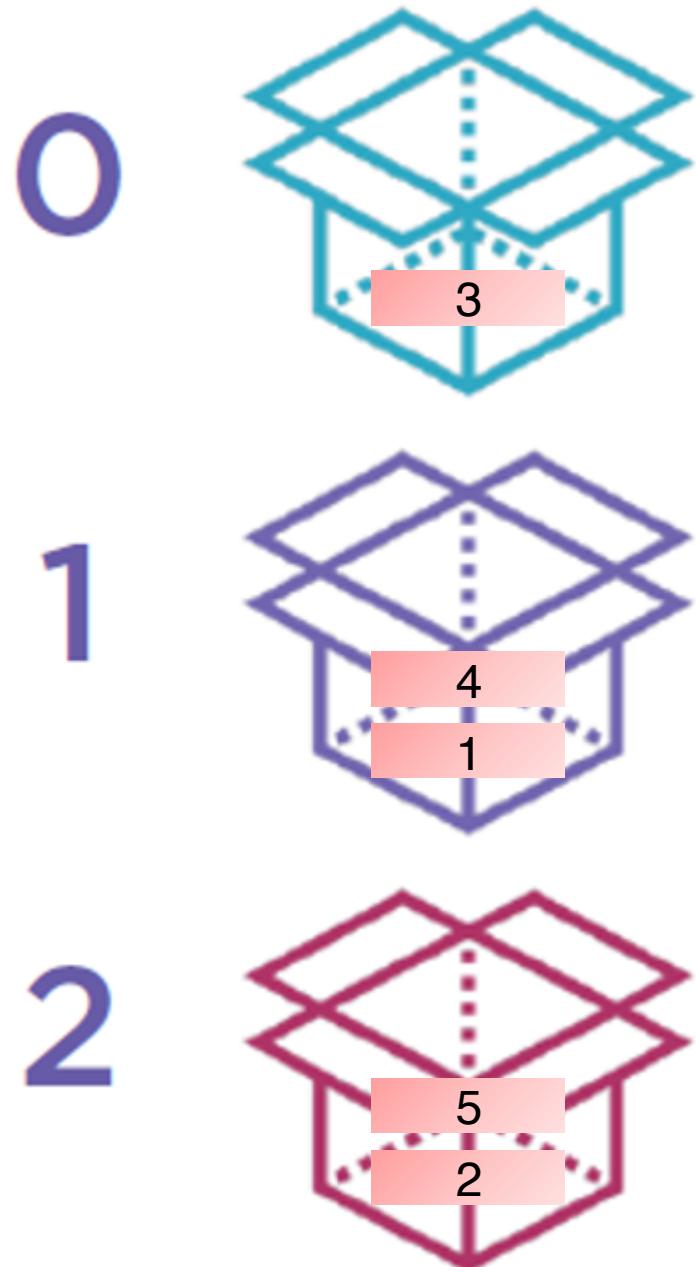


Bucketing

pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transac	cat_estado
1	SANTADER	K	1000	599	SP
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP
5	RAPPI	2	500000	199	RR
6	NUBANK	K	25000	88	ES
7	GOOGLE	D	400000	45	ES
8	FACEBOOK	D	500000	25	RR

5 % 3 =

2



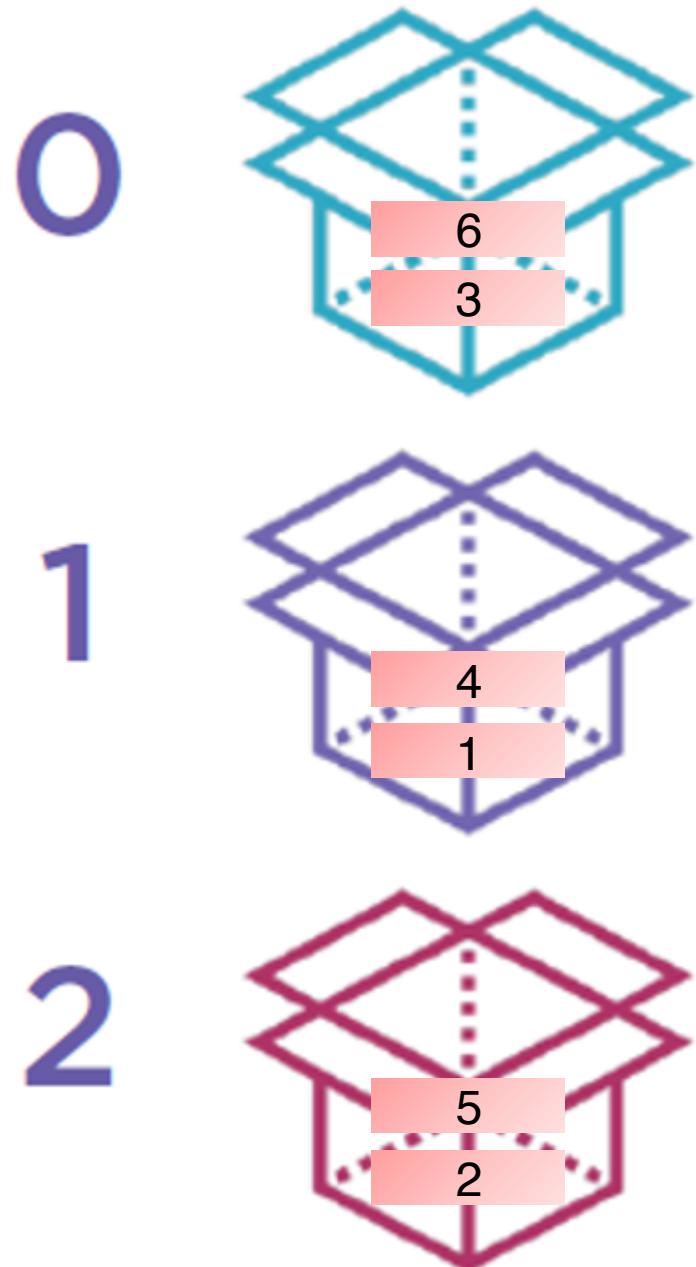


Bucketing

pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transac	cat_estado
1	SANTADER	K	1000	599	SP
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP
5	RAPPI	2	500000	199	RR
6	NUBANK	K	25000	88	ES
7	GOOGLE	D	400000	45	ES
8	FACEBOOK	D	500000	25	RR

6 % 3 =

0



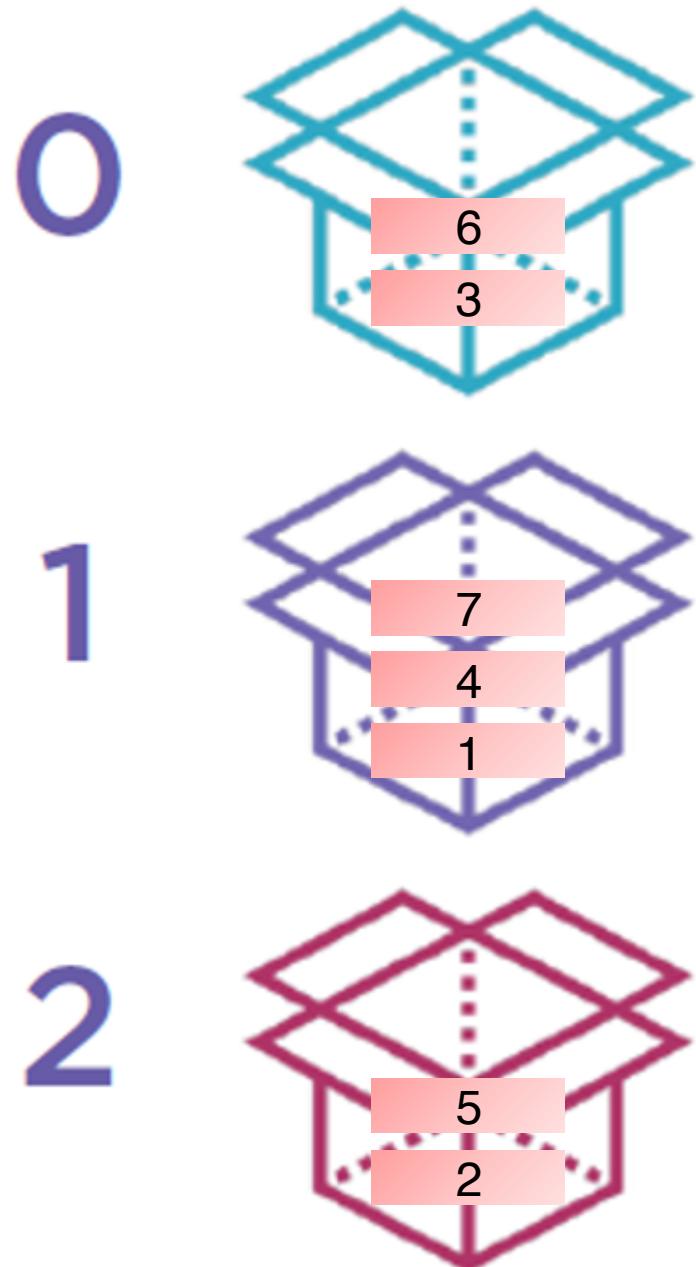


Bucketing

pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transac	cat_estado
1	SANTADER	K	1000	599	SP
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP
5	RAPPI	2	500000	199	RR
6	NUBANK	K	25000	88	ES
7	GOOGLE	D	400000	45	ES
8	FACEBOOK	D	500000	25	RR

7 % 3 =

1



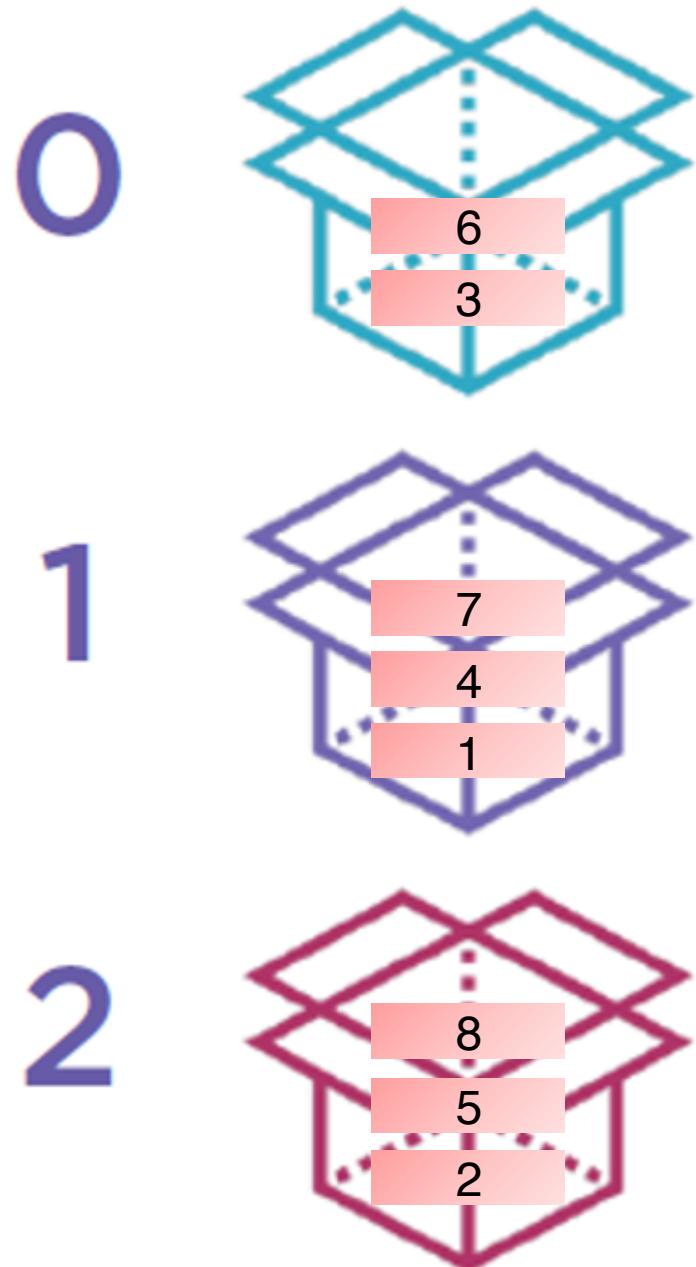


Bucketing

pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transac	cat_estado
1	SANTADER	K	1000	599	SP
2	UBER	K	25000	58	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP
5	RAPPI	2	500000	199	RR
6	NUBANK	K	25000	88	ES
7	GOOGLE	D	400000	45	ES
8	FACEBOOK	D	500000	25	RR

8 % 3 =

2





Bucketing

pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transac	cat_estado
1	SANTADER	K	1000	599	SP
2	UBER	K	25000	53	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP
5	RAPPI	2	500000	199	RR
6	NUBANK	K	25000	88	ES
7	GOOGLE	D	400000	45	ES
8	FACEBOOK	D	500000	25	RR



pk_cd_cnpj	str_name	fl_ever30
8	FACEBOOK	500000
7	GOOGLE	400000
6	NUBANK	25000
5	RAPPI	500000
4	IFOOD	100000
3	ITAU	400000
2	UBER	25000
1	SANTADER	1000

Precisa percorrer a tabela inteira



Bucketing

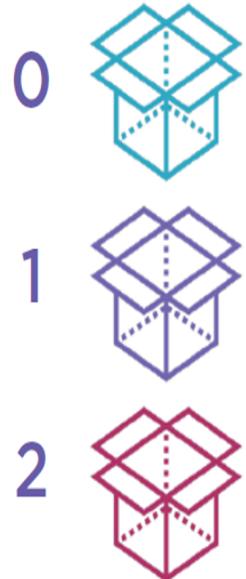
pk_cd_cnpj	str_name	cat_segmento	vlr_transacao	qtd_transac	cat_estado
1	SANTADER	K	1000	599	SP
2	UBER	K	25000	53	RJ
3	ITAU	2	400000	66	RJ
4	IFOOD	D	100000	99	SP
5	RAPPI	2	500000	199	RR
6	NUBANK	K	25000	88	ES
7	GOOGLE	D	400000	45	ES
8	FACEBOOK	D	500000	25	RR

pk_cd_cnpj	str_name	fl_ever30
6	NUBANK	25000
3	ITAU	400000

↓

pk_cd_cnpj	str_name	fl_ever30
7	GOOGLE	400000
4	IFOOD	100000
1	SANTADER	1000

pk_cd_cnpj	str_name	fl_ever30
8	FACEBOOK	500000
5	RAPPI	500000
2	UBER	25000



Precisa percorrer 1
bucket.



Window Functions

```
FUNCTION_NAME ( field )
OVER (window_clause) AS alias
```



Window Functions

**FUNCTION_NAME (field)
OVER (window_clause) AS alias**

firstname	department	startdate
Andrew	1	1/23/1999
Jacob	1	7/11/1990
Daniel	2	6/24/2004
Anna	1	10/7/2001
Pierre	1	2/22/2009
Ruth	2	6/6/1998
Anthony	1	11/29/1995
Isabella	2	9/28/1997
Jose	2	3/17/2013



Window Functions

```
SELECT firstname,  
       department,  
       startdate,  
       RANK() OVER (PARTITION BY  
department ORDER BY startdate ) AS rank  
FROM Employees;
```

firstname	department	startdate
Andrew	1	1/23/1999
Jacob	1	7/11/1990
Daniel	2	6/24/2004
Anna	1	10/7/2001
Pierre	1	2/22/2009
Ruth	2	6/6/1998
Anthony	1	11/29/1995
Isabella	2	9/28/1997
Jose	2	3/17/2013



Window Functions

PARTITION BY department

firstname	department	startdate
Andrew	1	1/23/1999
Jacob	1	7/11/1990
Daniel	2	6/24/2004
Anna	1	10/7/2001
Pierre	1	2/22/2009
Ruth	2	6/6/1998
Anthony	1	11/29/1995
Isabella	2	9/28/1997
Jose	2	3/17/2013

ORDER BY startdate

firstname	department	startdate
Andrew	1	1/23/1999
Jacob	1	7/11/1990
Anna	1	10/7/2001
Pierre	1	2/22/2009
Anthony	1	11/29/1995

RANK ()

firstname	department	startdate
Jacob	1	7/11/1990
Anthony	1	11/29/1995
Andrew	1	1/23/1999
Anna	1	10/7/2001
Pierre	1	2/22/2009

firstname	department	startdate	rank
Jacob	1	7/11/1990	1
Anthony	1	11/29/1995	2
Andrew	1	1/23/1999	3
Anna	1	10/7/2001	4
Pierre	1	2/22/2009	5



firstname	department	startdate
Ruth	2	6/6/1998
Daniel	2	6/24/2004
Jose	2	3/17/2013
Isabella	2	9/28/1997



firstname	department	startdate
Isabella	2	9/28/1997
Daniel	2	6/24/2004
Jose	2	3/17/2013
Ruth	2	6/6/2013

firstname	department	startdate	rank
Isabella	2	9/28/1997	1
Daniel	2	6/24/2004	2
Jose	2	3/17/2013	3
Ruth	2	6/6/2013	4



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER (ORDER  
BY date ROWS BETWEEN 2 PRECEDING  
AND CURRENT ROW) AS moving_avg  
FROM t_employee;
```

'80389123000101'	, 4500	, 201901
'80389123000103'	, 1500	, 201901
'80389123000104'	, 500	, 201901
'80389123000102'	, 4300	, 201901
'80389123000101'	, 4500	, 201902
'80389123000103'	, 1500	, 201902
'80389123000104'	, 500	, 201902
'80389123000102'	, 4300	, 201902
'80389123000101'	, 4500	, 201903
'80389123000103'	, 1500	, 201903
'80389123000104'	, 500	, 201903
'80389123000102'	, 4300	, 201903
'80389123000101'	, 4500	, 201904



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER (ORDER  
BY date ROWS BETWEEN 2 PRECEDING  
AND CURRENT ROW) AS moving_avg  
FROM t_employee;
```

id	vlr_transacao	date
'80389123000101'	4500	201901
'80389123000103'	1500	201901
'80389123000104'	500	201901
'80389123000102'	4300	201901
'80389123000101'	4500	201902
'80389123000103'	1500	201902
'80389123000104'	500	201902
'80389123000102'	4300	201902
'80389123000101'	4500	201903
'80389123000103'	1500	201903
'80389123000104'	500	201903
'80389123000102'	4300	201903
'80389123000101'	4500	201904



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER (ORDER  
BY date ROWS BETWEEN 2 PRECEDING  
AND CURRENT ROW) AS moving_avg  
FROM t_employee;
```

'80389123000101'	4500	201901
'80389123000103'	1500	201901
'80389123000104'	500	201901
'80389123000102'	4300	201901
'80389123000101'	4500	201902
'80389123000103'	1500	201902
'80389123000104'	500	201902
'80389123000102'	4300	201902
'80389123000101'	4500	201903
'80389123000103'	1500	201903
'80389123000104'	500	201903
'80389123000102'	4300	201903
'80389123000101'	4500	201904



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER (ORDER  
BY date ROWS BETWEEN 2 PRECEDING  
AND CURRENT ROW) AS moving_avg  
FROM t_employee;
```

'80389123000101'	4500	201901
'80389123000103'	1500	201901
'80389123000104'	500	201901
'80389123000102'	4300	201901
'80389123000101'	4500	201902
'80389123000103'	1500	201902
'80389123000104'	500	201902
'80389123000102'	4300	201902
'80389123000101'	4500	201903
'80389123000103'	1500	201903
'80389123000104'	500	201903
'80389123000102'	4300	201903
'80389123000101'	4500	201904



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER (ORDER  
BY date ROWS BETWEEN 2 PRECEDING  
AND CURRENT ROW) AS moving_avg  
FROM t_employee;
```

'80389123000101'	4500	201901)
'80389123000103'	1500	201901)
'80389123000104'	500	201901)
'80389123000102'	4300	201901)
'80389123000101'	4500	201902)
'80389123000103'	1500	201902)
'80389123000104'	500	201902)
'80389123000102'	4300	201902)
'80389123000101'	4500	201903)
'80389123000103'	1500	201903)
'80389123000104'	500	201903)
'80389123000102'	4300	201903)
'80389123000101'	4500	201904),



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER (ORDER  
BY date ROWS BETWEEN 2 PRECEDING  
AND CURRENT ROW) AS moving_avg  
FROM t_employee;
```

'80389123000101'	, 4500	, 201901)
'80389123000103'	, 1500	, 201901)
'80389123000104'	, 500	, 201901)
'80389123000102'	, 4300	, 201901)
'80389123000101'	, 4500	, 201902)
'80389123000103'	, 1500	, 201902)
'80389123000104'	, 500	, 201902)
'80389123000102'	, 4300	, 201902)
'80389123000101'	, 4500	, 201903)
'80389123000103'	, 1500	, 201903)
'80389123000104'	, 500	, 201903)
'80389123000102'	, 4300	, 201903)
'80389123000101'	, 4500	, 201904),



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER (ORDER  
BY date ROWS BETWEEN 2 PRECEDING  
AND CURRENT ROW) AS moving_avg  
FROM t_employee;
```

'80389123000101'	4500	201901
'80389123000103'	1500	201901
'80389123000104'	500	201901
'80389123000102'	4300	201901
'80389123000101'	4500	201902
'80389123000103'	1500	201902
'80389123000104'	500	201902
'80389123000102'	4300	201902
'80389123000101'	4500	201903
'80389123000103'	1500	201903
'80389123000104'	500	201903
'80389123000102'	4300	201903
'80389123000101'	4500	201904



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER (ORDER  
BY date ROWS BETWEEN 2 PRECEDING  
AND CURRENT ROW) AS moving_avg  
FROM t_employee;
```

'80389123000101'	4500	201901)
'80389123000103'	1500	201901)
'80389123000104'	500	201901)
'80389123000102'	4300	201901)
'80389123000101'	4500	201902)
'80389123000103'	1500	201902)
'80389123000104'	500	201902)
'80389123000102'	4300	201902)
'80389123000101'	4500	201903)
'80389123000103'	1500	201903)
'80389123000104'	500	201903)
'80389123000102'	4300	201903)
'80389123000101'	4500	201904),



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER (ORDER  
BY date ROWS BETWEEN 2 PRECEDING  
AND CURRENT ROW) AS moving_avg  
FROM t_employee;
```

'80389123000101'	4500	201901
'80389123000103'	1500	201901
'80389123000104'	500	201901
'80389123000102'	4300	201901
'80389123000101'	4500	201902
'80389123000103'	1500	201902
'80389123000104'	500	201902
'80389123000102'	4300	201902
'80389123000101'	4500	201903
'80389123000103'	1500	201903
'80389123000104'	500	201903
'80389123000102'	4300	201903
'80389123000101'	4500	201904



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER  
(PARTITION BY pk_cd_cnpj ORDER BY date  
ROWS BETWEEN 2 PRECEDING AND  
CURRENT ROW) AS moving_avg  
FROM t_employee;
```

('80389123000101',	4300,	201901)
('80389123000101',	4400,	201902)
('80389123000101',	4500,	201903)
('80389123000101',	4900,	201904)
('80389123000101',	4400,	201905)
('80389123000103',	1600,	201901)
('80389123000103',	1200,	201902)
('80389123000103',	1100,	201903)
('80389123000103',	1300,	201904)
('80389123000103',	1300,	201905)



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER  
(PARTITION BY pk_cd_cnpj ORDER BY date  
ROWS BETWEEN 2 PRECEDING AND  
CURRENT ROW) AS moving_avg  
FROM t_employee;
```

('80389123000101',	4300,	201901)
('80389123000101',	4400,	201902)
('80389123000101',	4500,	201903)
('80389123000101',	4900,	201904)
('80389123000101',	4400,	201905)
('80389123000103',	1600,	201901)
('80389123000103',	1200,	201902)
('80389123000103',	1100,	201903)
('80389123000103',	1300,	201904)
('80389123000103',	1300,	201905)



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER  
(PARTITION BY pk_cd_cnpj ORDER BY date  
ROWS BETWEEN 2 PRECEDING AND  
CURRENT ROW) AS moving_avg  
FROM t_employee;
```

('80389123000101', 4300, 201901)	
('80389123000101', 4400, 201902)	
('80389123000101', 4500, 201903)	
('80389123000101', 4900, 201904)	
('80389123000101', 4400, 201905)	
('80389123000103', 1600, 201901)	
('80389123000103', 1200, 201902)	
('80389123000103', 1100, 201903)	
('80389123000103', 1300, 201904)	
('80389123000103', 1300, 201905)	



Window Functions

```
SELECT *,  
       AVG(vlr_transacao) OVER  
(PARTITION BY pk_cd_cnpj ORDER BY date  
ROWS BETWEEN 2 PRECEDING AND  
CURRENT ROW) AS moving_avg  
FROM t_employee;
```

('80389123000101',	4300,	201901)
('80389123000101',	4400,	201902)
('80389123000101',	4500,	201903)
('80389123000101',	4900,	201904)
('80389123000101',	4400,	201905)
('80389123000103',	1600,	201901)
('80389123000103',	1200,	201902)
('80389123000103',	1100,	201903)
('80389123000103',	1300,	201904)
('80389123000103',	1300,	201905)